

# Math 128A, Summer 2019

Lecture 2, Tuesday 6/25/2019

## 1 Review of Lec 1 Floating Point Arithmetic

In last lecture, we ended with rounding FPN (Floating Point Numbers), and tie-breaking.

**Definition: Normal number -**

$$[S] \quad [C] \quad [f]$$

$$(-1)^s \cdot 2^{(C-1023)} \cdot (1 + f)$$

(65-bit effective by adding the implied 1 in mantissa)

**Definition: Subnormal number -**

$$C := 0$$

$$\implies (-1)^S \cdot 2^{-1022} \cdot (f)$$

“Gradual underflow”. We lose the  $1 + \dots$ , but we get another 51 bits before our numbers go to zero  $2^{-1023} \rightarrow 2^{-1022}$ .

Note: We used to use “fixed-point” arithmetic (with all non-sign bits assigned to the mantissa; i.e. no characteristic value) back when computers were vacuum tubes in a room. Now with the characteristic, we get better usability.

Other special cases:

Consider that  $\pm 0$  are both “standard numbers”. “ $\pm \text{Inf}$ ” as well. But there’s NaN (not a number), which is anything that has no valid representation.

### 1.1 Spacing of FPNs

**Definition: Machine Epsilon (mach eps) -**

Also called **ULP** for “unit in the last place”.

$$\varepsilon := 2^{-52}$$

So we have the following spacing:

$$1, 1 + \varepsilon, 1 + 2\varepsilon, \dots, 2 - \varepsilon, 2, 2 + 2\varepsilon, 2 + 4\varepsilon, \dots, 4 - 2\varepsilon, 4, \dots$$

**Example:** Constructing the IEEE FP representations:

$$\begin{aligned}
 1 &= S = 0, C = 0111111111 = 1023, f = 0 \cdots 00 \\
 2 &= S = 0, C = 1000000000 = 1024, f = 0 \cdots 00 \\
 2 + 2\varepsilon &= S = 0, C = 1000000000 = 1024, f = 0 \cdots 01 \\
 2 + 4\varepsilon &= S = 0, C = 1000000000 = 1024, f = 0 \cdots 10 \\
 2 + 6\varepsilon &= S = 0, C = 1000000000 = 1024, f = 0 \cdots 11 \\
 1 - \frac{\varepsilon}{2} &= S = 0, C = 1022, f = 1 \cdots 1
 \end{aligned}$$

**Definition: Rounding -**

Rounding:  $fl : \mathbb{R} \rightarrow$  nearest FPN

**Example:**

$$fl\left(1 + \varepsilon + \frac{\varepsilon}{\pi}\right) = 1 + \varepsilon fl\left(1 - \frac{\varepsilon}{\pi}\right)$$

**Resolving ties:** (to last bit zero)

$$\begin{aligned}
 fl\left(1 + \frac{\varepsilon}{2}\right) &= 1 \\
 fl\left(1 - \frac{\varepsilon}{2}\right) &= 1 - \frac{\varepsilon}{2} \\
 fl\left(1 - \frac{\varepsilon}{4}\right) &= 1 \quad (\text{this is more interesting})
 \end{aligned}$$

Note that a real number will live in between two FPNs, one will be even and one will be odd. It's not necessarily rounding down; just rounding to 0. This prevents bias accumulating ("drift") over many operations.

**Remark:** Rounding is "monotone", which means it preserves order (inequalities).

$$a < b \implies fl(a) \leq fl(b)$$

Note the possible equality if they both round to the same FPN.

**Aside/Motivation:** debugging a program to run on supercomputer

$$\arccos(x); x = \frac{a}{a^2 + b^2}$$

But now with the IEEE standard, we don't have these issues with inequalities, where values go out of 'bounds'.

## 2 Rounding Error

$$\begin{aligned}
 1 \leq x < 2 : \quad & |fl(x) - x| \leq \varepsilon \\
 & \leq \frac{\varepsilon}{2} \\
 & \leq \frac{\varepsilon}{2} |x| \\
 \text{Relative Error} =: & \frac{|fl(x) - x|}{|x|} \leq \frac{\varepsilon}{2} \leq \varepsilon
 \end{aligned}$$

But of course, if we have a large  $x$ , our precision is much smaller because we need to use bits to store the size.

If we have the following addition:

$$\begin{aligned}
 & \begin{array}{r}
 \overbrace{0.001 \cdots 1}^{52} \\
 + 1.000 \cdots 0 \\
 \hline
 = \underbrace{1.001 \cdots 1}_{54} 11 = 1.010 \cdots 0 \\
 = [2^{-3} \cdot (1 - \varepsilon)] + [1] = \underbrace{1 + 2^{-3} - 2^{-3}\varepsilon}_{\text{exactly}}
 \end{array} \\
 fl(x, y) &= fl(1 + 2^{-3} - 2^{-3}\varepsilon) = 1 + 2^{-3}
 \end{aligned}$$

**Mantra of FP Arithmetic:** Let  $x, y$  be FPN, floating point numbers. The floating point result of any binary operation on two FPNs gives us:

$$fl(x + y) = \text{the exact result, correctly rounded}$$

And likewise for  $fl(xy)$ ,  $fl(x - y)$ ,  $fl(x/y)$ ,  $fl(\sqrt{x})$ .

So instead of doing FP arithmetic per term, we just perform the regular math, then round at the end.

**Example:**  $0.00 \underbrace{11 \cdots 11}_{52} = 2^{-3} \times 1. \underbrace{1 \cdots 1}_{52}$

Or equivalently,  $fl(1 + 2^{-3}(2 - \varepsilon)) = fl(1 + 2^{-2} - 2^{-3}\varepsilon) = 1.01$ .  
So we have:

**Definition: Relative Error**  $|\delta| \leq \frac{\varepsilon}{2}$  -

$$\begin{aligned}
 \frac{|fl(x) - x|}{|x|} &\leq \frac{\varepsilon}{2} \\
 \underbrace{|fl(x) - x|}_{\delta x} &\leq \frac{\varepsilon}{2} |x| \\
 fl(x) - x &= \delta x \\
 fl(x) &= x + \delta x \\
 &= x(1 + \delta), \quad |\delta| \leq \frac{\varepsilon}{2}
 \end{aligned}$$

**Remark:**  $\delta$  is INDEPENDENT of  $x$ .

Thus:

$$\begin{aligned}
 fl(x \pm y) &= (x \pm y)(1 + \delta), & |\delta| &\leq \frac{\varepsilon}{2} \\
 fl(\sqrt{x}) &= \sqrt{x}(1 + \delta), & |\delta| &\leq \frac{\varepsilon}{2} \\
 fl((x + y) + z) &\neq (x + y + z)(1 + \delta) \\
 &= fl[fl(x + y) + z] \\
 &= fl[(x + y)(1 + \delta) + z] \\
 &= [(x + y)(1 + \delta_1) + z](1 + \delta_2) \\
 &= (x + y)(1 + \delta_1)(1 + \delta_2) + z(1 + \delta_2) \\
 &= (x + y)(1 + \delta_1 + \delta_2 + \delta_1\delta_2) + z(1 + \delta_2) \\
 &= (x + y)(1 + 2\delta_3) + z(1 + \delta_2), & |\delta_3| &\leq \frac{3}{2} \\
 &= x(1 + 2\delta_3) + y(1 + 2\delta_3) + z(1 + \delta_2)
 \end{aligned}$$

Note that these deltas can be different. But each of these terms are very close to their respective true values  $x, y, z$ .

**Definition: Backwards Error Analysis -**

Take  $fl(x + y + z) = \hat{x} + \hat{y} + \hat{z}$  where:

$$\begin{aligned}
 \hat{x} &:= x(1 + 2\delta_3) \\
 \hat{y} &:= y(1 + 2\delta_3) \\
 \hat{z} &:= z(1 + \delta_2)
 \end{aligned}$$

As opposed to...

**Definition: Forward Error Analysis -**

$$\begin{aligned}
 \frac{|fl(x + y + z) - (x + y + z)|}{|x + y + z|} &\leq \frac{|(x + y)(1 + 2\delta_3) + z(1 + \delta_2) - (x + y + z)|}{|x + y + z|} \\
 &\leq \frac{2|\delta_3||x + y| + |z||\delta_2|}{|x + y + z|} \\
 &\leq \frac{2|x + y| + |z|}{|x + y + z|} \frac{\varepsilon}{2}
 \end{aligned}$$

**Example:** Take  $x = 1 + \varepsilon, y = -\varepsilon, z = -1$ . Recall that we had  $1 \leq x < 2$ , with  $|fl(x) - x| \leq \frac{\varepsilon}{2}|x|$ .

Our relative error bound is

$$\frac{2|x + y| + |z|}{|x + y + z|} \frac{\varepsilon}{2} \leq \frac{3}{0} \cdot \frac{\varepsilon}{2}$$

But we have division by zero, which is useless, so our result would be unverifiable.

**Mantra 2 of Computing:** Don't compute 0 by subtracting nonzero objects (because then we can't tell the relative error).

**Key Takeaway:**

$$fl(x) = x(1 + \delta), \quad |\delta| \leq \frac{\varepsilon}{2}$$

**Example of large number of operations, backwards error analysis:**

Suppose we want to find:

$$S_n = \sum_{j=1}^n x_j$$

We need an algorithm:

$$S_1 = x_1 \tag{1}$$

$$S_2 = S_1 + x_2 \tag{2}$$

$$\vdots \tag{3}$$

$$S_n = S_{n-1} + x_n \tag{4}$$

Pseudocode

```

1 S_n = S_{n-1} + x_n
2 unless
3   n = 1    when S_1 = x_1
4 recursion

```

In numerical analysis we tend to avoid recursion because it can be expensive (number of operations).

We conclude:

$$\begin{aligned}
 fl(S_n) &= x_1[1 + (n-1)\delta_1] \\
 &\quad + x_2[1 + (n-1)\delta_2] \\
 &\quad + x_3[1 + (n-2)\delta_3] \\
 &\quad + \cdots + x_n[1 + \delta_n]
 \end{aligned}$$

or more sloppily (overestimating),

$$\begin{aligned}
 fl(S_n) &= x_1[1 + (n)\delta_1] \\
 &\quad + x_2[1 + (n-1)\delta_2] \\
 &\quad + x_3[1 + (n-2)\delta_3] \\
 &\quad + \cdots + x_n[1 + \delta_n]
 \end{aligned}$$

So BEA (backwards error analysis) says that  $x$  has  $n$  errors.

If we want FEA (forward error analysis) and its diagnostic, then we look at:

$$\left| \frac{fl(S_n) - S_n}{S_n} \right| \leq \frac{n|x_1| + (n-1)|x_2| + \cdots + 1|x_n|}{|S_n|} \cdot \frac{\varepsilon}{2}$$

**Remark:** Notes end here for today.

Big ideas: BEA, FEA,  $(1 + \delta_n)$  and simplifications for compounding error from a sequence of operations.